

Final project guidelines

Note: Use these guidelines if and only if you are pursuing a **final project of your own design**. For those taking the final exam instead of the project, see the (separate) final exam notebook.

The task

Your task is to: identify an interesting problem connected to the humanities or humanistic social sciences that's addressable with the help of computational methods, formulate a hypothesis about it, devise an experiment or experiments to test your hypothesis, present the results of your investigations, and discuss your findings.

These tasks essentially replicate the process of writing an academic paper. You can think of your project as a paper in miniature.

You are free to present each of these tasks as you see fit. You should use narrative text (that is, your own writing in a markdown cell), citations of others' work, numerical results, tables of data, and static and/or interactive visualizations as appropriate. Total length is flexible and depends on the number of people involved in the work, as well as the specific balance you strike between the ambition of your question and the sophistication of your methods. But be aware that numbers never, ever speak for themselves. Quantitative results presented without substantial discussion will not earn high marks.

Your project should reflect, at minimum, ten **or more** hours of work by each participant, though you will be graded on the quality of your work, not the amount of time it took you to produce it. Most high-quality projects represent twenty or more hours of work by each member.

Pick an important and interesting problem!

No amount of technical sophistication will overcome a fundamentally uninteresting problem at the core of your work. You have seen many pieces of successful computational humanities research over the course of the semester. You might use these as a guide to the kinds of problems that interest scholars in a range of humanities disciplines. You may also want to spend some time in the library, reading recent books and articles in the professional literature. **Problem selection and motivation are integral parts of the project.** Do not neglect them.

Format

You should submit your project as a PDF document created using the included *LATEX* template. Consult the template for information on formatting and what is expected in each section. You can use your favorite text editor or something like [Overleaf](#) to edit this document. You will also submit this Jupyter notebook, along with all data necessary to reproduce your analysis. If your dataset is too large to share easily, let us know in advance so that we can find a workaround.

All code used in the project should be present in the notebook (except for widely-available libraries that you import), but **be sure that we can read and understand your report in full without rerunning the code.**

Because you are submitting essentially a mini-paper in the PDF writeup, I don't have any particular formatting expectations for written material in this notebook. However, you should include **all code used when completing the final project, with comments added for clarity.** It should be straightforward to map code from the notebook to sections/figures/results in your paper, and vice versa.

Grading

This project takes the place of the take-home final exam for the course. It is worth 35% of your overall grade. You will be graded on

the quality and ambition of each aspect of the project. No single component is more important than the others.

Practical details

- The project is due at **4:30 PM EST on Wednesday, December 17** via upload to CMS of a single zip file containing your fully executed Jupyter notebook report, a PDF copy of the notebook, and all associated data. **You may not use slip days for the final project or exam.**
 - You may work alone or in a group of up to three total members.
 - If you work in a group, be sure to list the names of the group members.
 - For groups, create your group on CMS and submit one notebook for the entire group. **Each group should also submit a statement of responsibility** that describes in general terms who performed which parts of the project.
 - You may post questions on Ed, but should do so privately (visible to course staff only).
-

Your info

- NetID(s): kgc42
 - Name(s): Kyle Chu
-

```
In [1]: import pandas as pd
import numpy as np
import kagglehub
import os
import nltk
from nltk.sentiment import SentimentIntensityAnalyzer
import spacy

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.decomposition import LatentDirichletAllocation
import re
```

```
from scipy import stats
import matplotlib.pyplot as plt
from matplotlib.patches import Patch
import seaborn as sns

from sklearn.linear_model import LogisticRegression
```

Step 1: Data Loading & Filtering

We'll load the Steam Reviews dataset, filter to Rainbow Six Siege English reviews, and remove short reviews. Class balancing will be performed after sentiment analysis, once we have filtered to negative reviews only.

```
In [2]: path = kagglehub.dataset_download("najzekeo/steam-reviews-2021")
print("Path to dataset files:", path)

GAME_APPID = 359550
```

Path to dataset files: /Users/kych2204/.cache/kagglehub/datasets/najzekeo/steam-reviews-2021/versions/1

Load in the data

```
In [3]: # load the CSV file

csv_path = os.path.join(path, "steam_reviews.csv")
print(f"Loading data from {csv_path}...")

chunk_size = 50000
chunks = []
total_read = 0
chunks_processed = 0

for chunk in pd.read_csv(csv_path, chunksize=chunk_size, low_
    chunks_processed += 1

    chunk['app_id'] = pd.to_numeric(chunk['app_id'], errors=_

        chunk_filtered = chunk[
            (chunk['app_id'] == GAME_APPID) &
            (chunk['language'].str.lower() == 'english')
        ].copy()
```

```
if len(chunk_filtered) > 0:  
    chunks.append(chunk_filtered)  
    total_read += len(chunk_filtered)  
  
df = pd.concat(chunks, ignore_index=True)  
print(f"\nTotal Rainbow Six Siege English reviews loaded: {len(df)}")
```

Loading data from /Users/kych2204/.cache/kagglehub/datasets/najzeke/steam-reviews-2021/versions/1/steam_reviews.csv...

Total Rainbow Six Siege English reviews loaded: 371,154

Clean and filter the data

```
In [4]: df['recommended'] = df['recommended'].astype(bool)  
  
# remove null/empty reviews  
initial_count = len(df)  
df = df[df['review'].notna()].copy()  
df = df[df['review'].str.strip() != ''].copy()  
print(f"After removing null/empty reviews: {len(df)} (removing {initial_count - len(df)})")  
  
# only keep reviews from users who purchased the game  
before_purchase_filter = len(df)  
df = df[df['steam_purchase'] == True].copy()  
print(f"After filtering to steam_purchase == True: {len(df)} (removing {before_purchase_filter - len(df)})")  
  
# remove very short reviews under 10 words  
df['token_count'] = df['review'].str.split().str.len()  
df = df[df['token_count'] >= 10].copy()  
print(f"After removing reviews with < 10 tokens: {len(df)}")  
  
# check class distribution  
print(f"\nClass distribution:")  
print(df['recommended'].value_counts())  
  
df.head()
```

```
After removing null/empty reviews: 370,156 (removed 998)
After filtering to steam_purchase == True: 312,169 (removed 5
7,987)
After removing reviews with < 10 tokens: 108,264
```

Class distribution:

recommended

True 90053

False 18211

Name: count, dtype: int64

Out [4]:

	Unnamed: 0	app_id	app_name	review_id	language	review
3	5860060	359550	Tom Clancy's Rainbow Six Siege	84794099	english	servers are shit actual shit but the game itse...
5	5860063	359550	Tom Clancy's Rainbow Six Siege	84793878	english	i do recomend this game bc it is rlly fun anf ...
9	5860080	359550	Tom Clancy's Rainbow Six Siege	84792481	english	This game is just awesome, the graphics and ev...
11	5860083	359550	Tom Clancy's Rainbow Six Siege	84792431	english	this game drives me crazy, but im still playin...
19	5860098	359550	Tom Clancy's Rainbow Six Siege	84791184	english	A good game overall. The only downside that I ...

5 rows × 24 columns

In [5]:

```
# remove unneeded columns
keep_cols = ['review', 'recommended', 'app_id', 'language',
             df = df[keep_cols].copy()

print(f"\nDataset shape: {df.shape}")

df.head()
```

Dataset shape: (108264, 6)

Out[5]:

		review	recommended	app_id	language	token_count	times
		servers are shit actual shit but the game itse...					
3			True	359550	english		44
5		i do recomend this game bc it is rlly fun anf ...	True	359550	english		16
9		This game is just awesome, the graphics and ev...	True	359550	english		117
11		this game drives me crazy, but im still playin...	False	359550	english		18
19		A good game overall. The only downside that I ...	True	359550	english		67

Step 2: Sentiment Analysis & Negative Subset

We'll use VADER sentiment analysis to identify negative reviews, then filter to only negative reviews for our analysis.

In [6]:

```
try:  
    nltk.data.find('vader_lexicon')
```

```
except LookupError:  
    nltk.download('vader_lexicon')  
  
sia = SentimentIntensityAnalyzer()  
  
[nltk_data] Downloading package vader_lexicon to  
[nltk_data]     /Users/kych2204/nltk_data...  
[nltk_data]     Package vader_lexicon is already up-to-date!
```

In [7]: # Configuration for sentiment analysis
NEG_THRESHOLD = -0.2

In [8]: # Compute sentiment scores for each review
print("Computing sentiment scores...")
df['sentiment_compound'] = df['review'].apply(lambda x: sia.

print(f"Sentiment score statistics:")
print(f" Mean: {df['sentiment_compound'].mean():.3f}")
print(f" Median: {df['sentiment_compound'].median():.3f}")
print(f" Min: {df['sentiment_compound'].min():.3f}")
print(f" Max: {df['sentiment_compound'].max():.3f}")

Computing sentiment scores...

Sentiment score statistics:
Mean: 0.394
Median: 0.637
Min: -1.000
Max: 1.000

In [9]: # Filter to negative reviews only
print(f"\nFiltering to negative reviews (sentiment_compound < 0)")
df_negative = df[df['sentiment_compound'] < NEG_THRESHOLD].copy()

print(f"Total negative reviews: {len(df_negative)}")
print(f"Removed {len(df) - len(df_negative)} non-negative reviews")

print(f"\nClass distribution in negative reviews:")
print(df_negative['recommended'].value_counts())

Filtering to negative reviews (sentiment_compound < -0.2)...

Total negative reviews: 20,910

Removed 87,354 non-negative reviews

Class distribution in negative reviews:

recommended

True 11461

False 9449

Name: count, dtype: int64

In [10]: # balance datasets

```
negative_recommended = df_negative[df_negative['recommended']]
negative_not_recommended = df_negative[df_negative['recommended'] == False]

print(f"\nNegative + Recommended before balancing: {len(negative_recommended)}")
print(f"Negative + Not Recommended before balancing: {len(negative_not_recommended)}")

n_sample = min(len(negative_recommended), len(negative_not_recommended))
print(f"\nSampling {n_sample} reviews from each class...")

negative_rec_sample = negative_recommended.sample(n=n_sample)
negative_not_rec_sample = negative_not_recommended.sample(n=n_sample)

df_negative_balanced = pd.concat([negative_rec_sample, negative_not_rec_sample])
df_negative_balanced = df_negative_balanced.sample(frac=1, random_state=42)

print(f"\nFinal balanced negative dataset: {len(df_negative_balanced)} reviews")
print("Class distribution after balancing:")
print(df_negative_balanced['recommended'].value_counts())
```

Negative + Recommended before balancing: 11,461

Negative + Not Recommended before balancing: 9,449

Sampling 9,449 reviews from each class...

Final balanced negative dataset: 18,898 reviews

Class distribution after balancing:

recommended

False 9449

True 9449

Name: count, dtype: int64

In [11]: # display 10 random reviews from balanced dataset of negative reviews

```
print("Sample of 10 negative reviews (balanced dataset):")
print("=" * 80)
```

```
sample_reviews = df_negative_balanced.sample(10)

for idx, row in sample_reviews.iterrows():
    rec_status = "RECOMMENDED" if row['recommended'] else "NOT RECOMMENDED"
    sentiment = row['sentiment_compound']
    print(f"\n[{rec_status}] (Sentiment: {sentiment:.3f})")
    print(f"Review: {row['review']}")
    print("-" * 80)
```

Sample of 10 negative reviews (balanced dataset):

[RECOMMENDED] (Sentiment: -0.930)

Review: this is not an easy FPS game for newbie, if you have experiences about FPS game, try it. If not, dont buy it because you will be very angry and dont know why you die and how you die. Anyway, it is a hardcore game.

[NOT RECOMMENDED] (Sentiment: -0.859)

Review: This game is the worst tryhard trash ive ever seen. I hope every developer and every player of this game will get b onecancer and their death will be fuckin painfully

[NOT RECOMMENDED] (Sentiment: -0.426)

Review: Its a great game but the only problem I have it that people get matched up with people who are a higher level then you and your teammates, I see people who are 4x higher then my level and most of the time my team loses, not because they're bad, its because the higher level people has more skills then we low level do.

Over all I would rate this game a 6.5/10

[NOT RECOMMENDED] (Sentiment: -0.947)

Review: The game would be so much better if all the players would have some actual integrity and sportsmanship. This is probably the most toxic community I've ever been into. There are players who will straight up tell you to uninstall because they killed you once, without any actual reason to be toxic. Then they will proceed to comment on everything you do and always come with the same lame argument: "i'M n0t eVeN tRyInG". What's hilarious is that these players think and pretend to be good at the game by playing the same basic operators which require just 2 brain cells to use, operators which require no skill, I'm talking about Jager, Bandit, Ash etc. Some of them have straight up the worst Ranked Stats, I'm talking about 0.5/0.6 win rate and k/d. Their only argument and the most predictable one with which they try to hit back is: "hAhA y0u g0t

KiLlEd bY s0mEoNe w0rSe tHeN y0u". Is that supposed to offensive? Just because your stats are bad doesn't mean killing somebody makes him/her bad at the game. Anybody can kill somebody in a video game, it's not written in stone that good players can't get killed by bad players. Important note: this happens in casual. Yes, all of this happens in casual. It's so sad and pathetic seeing people do this just so they can feel above a random person they just found online through a video game from which nobody gains nothing (obviously). I can't even imagine how many insecurities these neckbeards have. Maybe I'm too normal to understand them because I play games for fun, but I guess I picked the wrong game.

Leaving the biggest downside of the game aside, I've also encountered multiple times some bugs regarding the sound. In one game I simply could not hear explosions, gunfire or footsteps, just the interface of the game when pressing "Esc" and some effects.

The only fun part of this game is creating a private game and playing just with friends, the rest of the game is usually stressful and making me more and more allergic to stupidity.

[NOT RECOMMENDED] (Sentiment: -0.587)

Review: The worst game i have ever played in the history of my gaming career. Would not recommend.

[RECOMMENDED] (Sentiment: -0.226)

Review: in the beginning it wasn't good enough. Then it was changed. many operators were added and the game play has been improved as well. BUT, the toxicity and racism of the community (North American Servers at least) hasn't been addressed almost at all. Hopefully the environment will soon be unsafe for such behaviors.

[RECOMMENDED] (Sentiment: -0.727)

Review: One of the best 1st person shooters in a long time in my opinion. I mean who doesn't like to tear down walls and kill the enemies on the other side!

[RECOMMENDED] (Sentiment: -0.689)

Review: if u want a fucking jage to fucking spawn kill u this
is a game for u

[RECOMMENDED] (Sentiment: -0.421)

Review: this is the game ive been looking to waste my life an
d money on

[NOT RECOMMENDED] (Sentiment: -0.258)

Review: A 5 year old game and I can't play more than 3 matches
without game breaking bugs. I went completely deaf my last
game.

```
In [12]: # turn recommended column into binary
df_negative_balanced['recommended'] = df_negative_balanced['recomme
df = df_negative_balanced.copy()

print(f"\nFinal dataset for analysis: {len(df)} negative reviews")
print(f"Columns: {df.columns.tolist()}")
df.head()
```

Final dataset for analysis: 18,898 negative reviews
Columns: ['review', 'recommended', 'app_id', 'language', 'token_count', 'timestamp_created', 'sentiment_compound']

	review	recommended	app_id	language	token_count	timestamp
0	Fucking shit game, a disgrace to Tom Clancy's ...		0 359550	english		14
1	Honestly filled with annoying ass trolls and s...		0 359550	english		33
2	I have so many hours into this broken game. Sa...	1	359550	english		17
3	with 1k hours, i can now say, that i fucking h...	1	359550	english		21
4	this game is full of bugs, its so fuckin shit ...	0	359550	english		38

Step 3: Text Preprocessing & Feature Engineering

We'll preprocess the text, create lexicon-based features for bugs/performance, monetization/price, and gameplay/fun, and fit a topic model.

Text Preprocessing

```
In [13]: def preprocess_text(text):
    '''Clean text for feature extraction'''
    text = text.lower()
    text = re.sub(r'http\S+|www\S+', '', text)
    return text

df['clean_text'] = df['review'].apply(preprocess_text)
print(f"Preprocessed {len(df)} reviews")

nlp = spacy.load("en_core_web_sm")

def lemmatize_text(text):
    doc = nlp(text)
    lemmas = [token.lemma_ for token in doc if token.is_alpha]
    return " ".join(lemmas)

df['clean_text'] = df['clean_text'].apply(lemmatize_text)
print(f"Lemmatized {len(df)} reviews for topic modeling")

df.head()
```

Preprocessed 18,898 reviews

Lemmatized 18,898 reviews for topic modeling

	review	recommended	app_id	language	token_count	timestamp
0	Fucking shit game, a disgrace to Tom Clancy's ...		0 359550	english		14
1	Honestly filled with annoying ass trolls and s...		0 359550	english		33
2	I have so many hours into this broken game. Sa...	1	359550	english		17
3	with 1k hours, i can now say, that i fucking h...	1	359550	english		21
4	this game is full of bugs, its so fuckin shit ...	0	359550	english		38

Lexicon-Based Features

In [14]: # Bug / Performance complaints

```
QUALITY_LEXICON = [
    "crash",
    "lag", "laggy",
    "fps",
    "bug",
```

```
"glitch",
"freeze",
"stutter",
"performance",
"optimization",
"optimize",
"broken",
"unplayable",
"frame",
"server",
"disconnect",
"ping",
"latency",
"drop",
]

# Monetization / Price complaints
MONETIZATION_LEXICON = [
    "microtransaction",
    "mtx",
    "lootbox",
    "dlc",
    "pay2win",
    "pay-to-win",
    "overpriced",
    "refund",
    "sale",
    "expensive",
    "price",
    "cost",
    "money",
    "greedy",
    "greed",
    "monetization",
    "premium",
    "skin",
    "cosmetic",
]

# Gameplay / Fun / Content
GAMEPLAY_LEXICON = [
    "fun",
    "gameplay",
    "combat",
    "balance",
    "content",
    "story",
```

```
[ "graphic",  
  "soundtrack",  
  "music",  
  "replay",  
  "grind",  
  "coop",  
  "friend",  
  "enjoyable",  
  "enjoy",  
  "addictive",  
  "satisfying",  
  "satisfaction",  
  "mechanic",  
  "skill",  
  "competitive",  
  "rank",  
  "matchmaking",  
  "strategy",  
  "tactical",  
  "teamwork",  
  "operator",  
  "map",  
  "weapon",  
 ]
```

	review	recommended	app_id	language	token_count	timestamp
0	Fucking shit game, a disgrace to Tom Clancy's ...		0 359550	english		14
1	Honestly filled with annoying ass trolls and s...		0 359550	english		33
2	I have so many hours into this broken game. Sa...	1	359550	english		17
3	with 1k hours, i can now say, that i fucking h...	1	359550	english		21
4	this game is full of bugs, its so fuckin shit ...	0	359550	english		38

Topic Modeling

```
In [16]: # Vectorize lemmatized text for topic modeling
vectorizer = CountVectorizer(
    min_df=5,
    max_df=0.8,
    max_features=5000,
    lowercase=False,
```

```
        stop_words='english',
        token_pattern=r'\b[a-z]+\b'
    )

X_counts = vectorizer.fit_transform(df['clean_text'])
print(f"Vocabulary size: {len(vectorizer.get_feature_names_out())}")
print(f"Document-term matrix shape: {X_counts.shape}")
```

Vocabulary size: 5,000
Document-term matrix shape: (18898, 5000)

In [17]: N_TOPICS = 7

```
print(f"\nFitting LDA model with {N_TOPICS} topics...")
lda = LatentDirichletAllocation(
    n_components=N_TOPICS,
    random_state=42,
    max_iter=20,
    learning_method='batch'
)

topic_model = lda.fit(X_counts)
print("Topic model fitted")

topic_proportions = lda.transform(X_counts)
print(f"Topic proportions shape: {topic_proportions.shape}")
```

Fitting LDA model with 7 topics...
Topic model fitted
Topic proportions shape: (18898, 7)

```
# Add topic proportion features to dataframe
for i in range(N_TOPICS):
    df[f'topic_{i}'] = topic_proportions[:, i]

print(f"Added {N_TOPICS} topic proportion features")

df.head()
```

Added 7 topic proportion features

Out[19]:

	review	recommended	app_id	language	token_count	timestamp
0	Fucking shit game, a disgrace to Tom Clancy's ...		0 359550	english		14
1	Honestly filled with annoying ass trolls and s...		0 359550	english		33
2	I have so many hours into this broken game. Sa...		1 359550	english		17
3	with 1k hours, i can now say, that i fucking h...		1 359550	english		21
4	this game is full of bugs, its so fuckin shit ...		0 359550	english		38

In [20]:

```
feature_names = vectorizer.get_feature_names_out()

print(f"\nTop 10 words for each of {N_TOPICS} topics:")
print("=" * 80)
for topic_idx in range(N_TOPICS):
    top_words_idx = topic_model.components_[topic_idx].argsort()
    top_words = [feature_names[i] for i in top_words_idx]
    top_weights = [topic_model.components_[topic_idx][i] for
```

```
print(f"\nTopic {topic_idx}:")
for word, weight in zip(top_words, top_weights):
    print(f"  {word:15s} {weight:.4f}")
```

Top 10 words for each of 7 topics:**Topic 0:**

game	3337.8344
buy	2023.3807
operator	1419.8778
edition	897.1421
money	773.8313
starter	732.1420
pay	707.1234
like	673.0530
just	551.2209
rainbow	513.9176

Topic 1:

shoot	2353.2717
kill	1962.0678
wall	1072.7551
enemy	924.1449
die	868.8772
head	729.5928
team	555.7396
hostage	511.7692
bullet	475.0468
headshot	451.3103

Topic 2:

ban	3162.3184
game	1461.7601
say	1252.4946
ubisoft	1174.7551
people	876.4793
word	743.1414
just	639.8539
chat	582.4904
toxic	527.6792
make	496.6211

Topic 3:

game	9238.1728
play	3595.9938
bad	1730.3633
good	1519.4262
just	1233.5854
like	1184.6457

fun	1143.1626
hate	1124.1983
hour	1082.8317
time	921.4525

Topic 4:

shit	4490.1411
game	4025.1874
fuck	1775.5565
server	1444.2049
broken	1435.9621
fucking	1195.8517
u	1055.7627
ubisoft	871.5515
rank	764.0519
fix	742.2058

Topic 5:

game	10145.2935
bad	1935.3863
gun	1865.0833
play	1690.8638
just	1589.6303
bug	1573.4294
fix	1517.2368
make	1388.8875
issue	1333.4542
new	1294.9800

Topic 6:

game	6217.3254
team	3128.1793
play	2590.4302
player	2276.8398
kill	1972.7674
match	1584.0719
time	1411.1408
people	1197.6837
like	1176.7799
just	1157.0121

```
In [21]: print(f"\nFinal dataset shape: {df.shape}")
feature_cols = ['quality_count', 'money_count', 'gamemplay_c
topic_cols = [f'topic_{i}' for i in range(N_TOPICS)]
all_features = feature_cols + topic_cols
```

```
print(f"Total number of features: {len(all_features)}")
print(f"Total features:{all_features}")
```

Final dataset shape: (18898, 20)
 Total number of features: 13
 Total features:['quality_count', 'money_count', 'gamenplay_count',
 'token_count', 'exclamation_count', 'profanity_count',
 'topic_0', 'topic_1', 'topic_2', 'topic_3', 'topic_4', 'topic_5',
 'topic_6']

Step 4: Descriptive Analysis

We'll compare lexicon features and topic proportions between negative reviews that still recommend the game and those that do not. This gives an initial, lecture-style view of whether gameplay, bugs/performance, and monetization talk differ across the two groups before fitting a formal regression model.

```
In [22]: print("Group sizes (negative reviews only):")
print(df['recommended'].value_counts().rename(index={0: 'Not
lex_cols = ['quality_count', 'money_count', 'gameplay_count']
summary_lex = df.groupby('recommended')[lex_cols].agg(['mean
summary_lex
```

Group sizes (negative reviews only):
 recommended
 Not recommended 9449
 Recommended 9449
 Name: count, dtype: int64

	quality_count			money_count		
	mean	std	median	mean	std	median
recommended						
0	1.223079	2.914038		0.0	0.493389	1.418120
1	0.710627	2.128585		0.0	0.210845	1.026913

```
In [23]: # t-test and p-value significance results
rows = []
for col in lex_cols:
    g_rec = df.loc[df['recommended'] == 1, col]
```

```

g_not = df.loc[df['recommended'] == 0, col]
t_stat, p_val = stats.ttest_ind(g_rec, g_not, equal_var=False)
rows.append({
    'feature': col,
    'mean_rec': g_rec.mean(),
    'mean_not': g_not.mean(),
    't_stat': t_stat,
    'p_value': p_val,
})

lex_test_results = pd.DataFrame(rows)
print("\nT-test results for lexicon features:")
lex_test_results

```

T-test results for lexicon features:

	feature	mean_rec	mean_not	t_stat	p_value
0	quality_count	0.710627	1.223079	-13.803817	4.091839e-43
1	money_count	0.210845	0.493389	-15.686327	4.513015e-55
2	gameplay_count	1.688273	1.421307	6.835429	8.440462e-12

```

In [24]: # graph for t-test and p-value significance results
try:
    plt.style.use('seaborn-v0_8-whitegrid')
except:
    try:
        plt.style.use('seaborn-whitegrid')
    except:
        plt.style.use('seaborn')
sns.set_palette("husl")

plt.rcParams['mathtext.default'] = 'regular'

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

lex_test_results['feature_label'] = lex_test_results['feature']
lex_test_results['feature_label'] = lex_test_results['feature'].map({
    'Quality': 'Quality/Bug',
    'Money': 'Monetization',
    'Gameplay': 'Gameplay'
})

```

```

lex_test_results_sorted = lex_test_results.copy()
lex_test_results_sorted['is_positive'] = lex_test_results_so
lex_test_results_sorted['abs_t_stat'] = lex_test_results_so
lex_test_results_sorted = lex_test_results_sorted.sort_values
lex_test_results_sorted = lex_test_results_sorted.drop(['is_')

colors = ['#2E86AB' if x < 0 else '#A23B72' for x in lex_te
bars = ax1.barh(lex_test_results_sorted['feature_label'], le

for i, (idx, row) in enumerate(lex_test_results_sorted.iterrows()):
    p_val = row['p_value']
    t_stat = row['t_stat']

    if p_val < 0.001:
        sig_text = '***'
    elif p_val < 0.01:
        sig_text = '**'
    elif p_val < 0.05:
        sig_text = '*'
    else:
        sig_text = ''

    x_pos = t_stat + (0.5 if t_stat > 0 else -0.5)
    ax1.text(x_pos, i, sig_text, ha='center' if t_stat > 0 e
                           va='center', fontsize=12, fontweight='bold')

ax1.axvline(x=0, color='black', linestyle='-', linewidth=0.8)

ax1.set_xlabel('T-statistic', fontsize=12, fontweight='bold')
ax1.set_ylabel('Feature', fontsize=12, fontweight='bold')
ax1.set_title('T-test Statistics by Feature\n(n(Recommended vs
ax1.grid(axis='x', alpha=0.3, linestyle='--')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)

legend_elements = [
    Patch(facecolor='#2E86AB', alpha=0.8, label='Lower in Rec
    Patch(facecolor='#A23B72', alpha=0.8, label='Higher in Re
    plt.Line2D([0], [0], marker='', linestyle='', label='***'
]
ax1.legend(handles=legend_elements, loc='upper right', fontst

lex_test_results_sorted['neg_log10_p'] = -np.log10(lex_test_
colors2 = ['#2E86AB' if x < 0 else '#A23B72' for x in lex_te

bars2 = ax2.barh(lex_test_results_sorted['feature_label'], le
                           color=colors2, alpha=0.8, edgecolor='black'

```

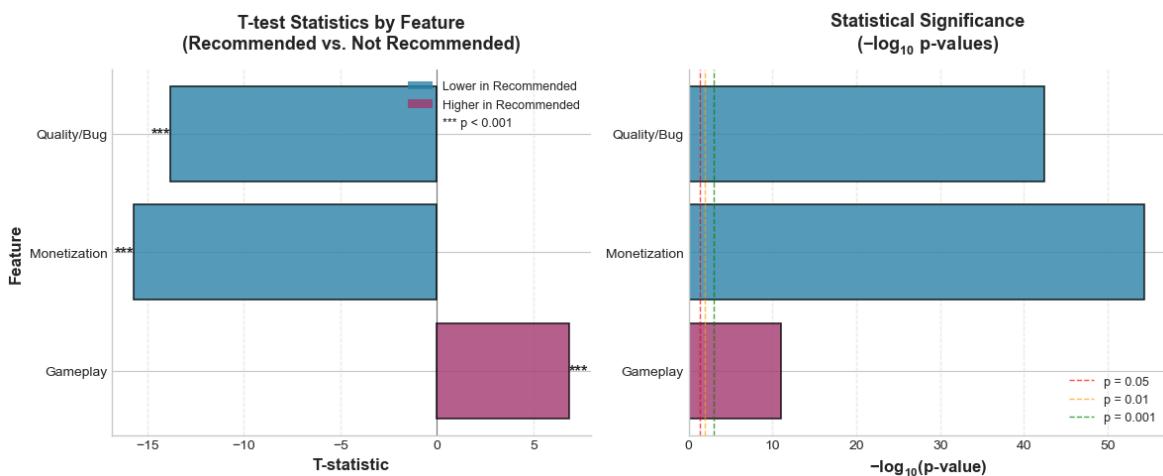
```

ax2.axvline(x=-np.log10(0.05), color='red', linestyle='--',
ax2.axvline(x=-np.log10(0.01), color='orange', linestyle='--',
ax2.axvline(x=-np.log10(0.001), color='green', linestyle='--')

ax2.set_xlabel(r'$-\log_{10}(p\text{-value})$', fontsize=12, fontweight='bold')
ax2.set_ylabel('')
ax2.set_title('Statistical Significance\n' + r'($-\log_{10}(p\text{-value})$)')
ax2.grid(axis='x', alpha=0.3, linestyle='--')
ax2.spines['top'].set_visible(False)
ax2.spines['right'].set_visible(False)
ax2.legend(loc='lower right', fontsize=9, framealpha=0.9)

plt.tight_layout()
plt.savefig('plots/t_test_results.png', dpi=300, bbox_inches='tight')
plt.show()

```



```

In [25]: df_all_neg = df_negative.copy()

df_all_neg['clean_text'] = df_all_neg['review'].apply(preprocess)
df_all_neg['clean_text'] = df_all_neg['clean_text'].apply(lowercase)

def count_lexicon_words(text, lexicon):
    tokens = text.split()
    return sum(1 for token in tokens if token in lexicon)

df_all_neg['quality_count'] = (
    df_all_neg['clean_text'].apply(lambda x: count_lexicon_words(x, quality_lexicon))
    / df_all_neg['token_count'] * 100
)
df_all_neg['money_count'] = (
    df_all_neg['clean_text'].apply(lambda x: count_lexicon_words(x, money_lexicon))
    / df_all_neg['token_count'] * 100
)

```

```

df_all_neg['gameplay_count'] = (
    df_all_neg['clean_text'].apply(lambda x: count_lexicon_words(x))
    / df_all_neg['token_count'] * 100
)

df_all_neg[['quality_count', 'money_count', 'gameplay_count']] = df_all_neg[['quality_count', 'money_count', 'gameplay_count']].fillna(0)

# Define groups
groups = {
    "All negative (full original dataset)": df_all_neg,
    "Negative & recommend": df_all_neg[df_all_neg['recommendation'] == 1],
    "Negative & not recommend": df_all_neg[df_all_neg['recommendation'] == 0]
}

rows = []
for gname, gdf in groups.items():
    rows.append({
        "group": gname,
        "Quality / Bugs": gdf['quality_count'].mean(),
        "Monetization": gdf['money_count'].mean(),
        "Gameplay / Fun": gdf['gameplay_count'].mean(),
    })

split_df = pd.DataFrame(rows)

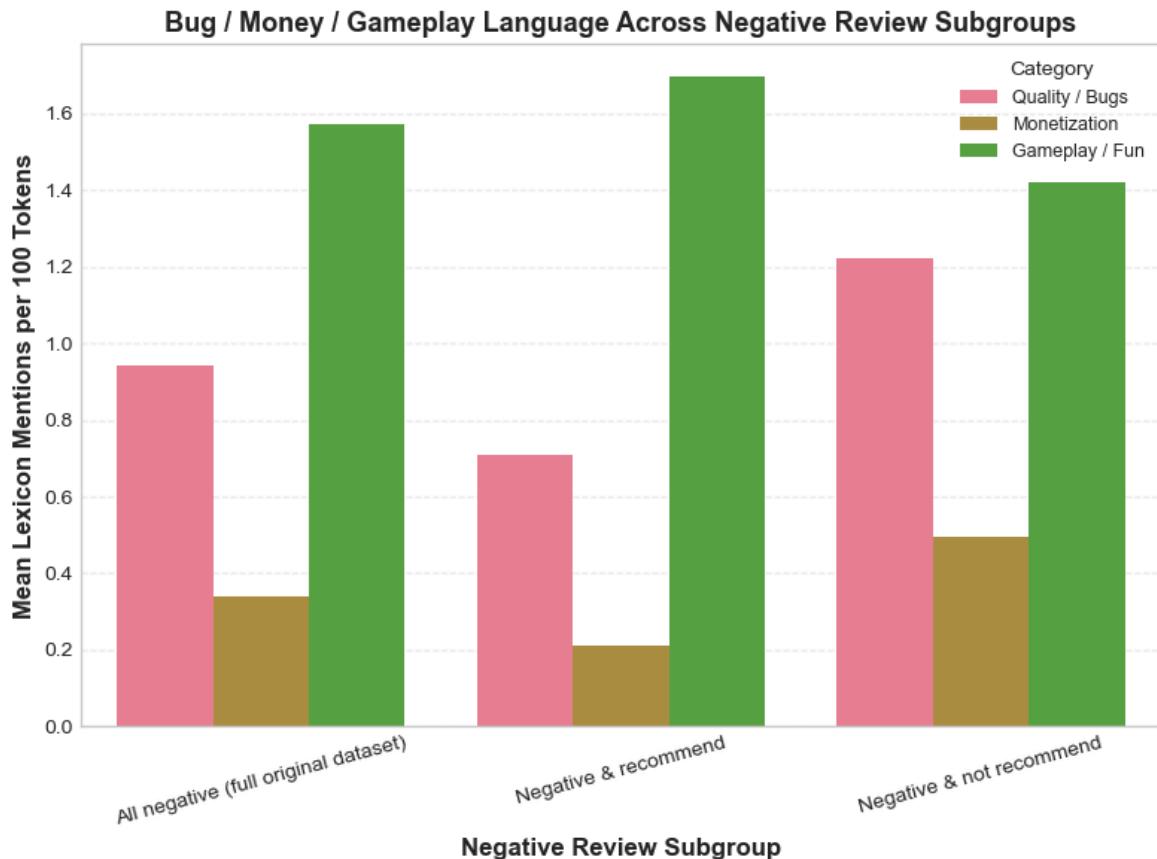
split_long = split_df.melt(
    id_vars='group',
    value_vars=['Quality / Bugs', 'Monetization', 'Gameplay / Fun'],
    var_name='Category',
    value_name='Percent_per_100_tokens'
)

plt.figure(figsize=(8, 6))
sns.barplot(
    data=split_long,
    x='group',
    y='Percent_per_100_tokens',
    hue='Category'
)

plt.xlabel('Negative Review Subgroup', fontsize=12, fontweight='bold')
plt.ylabel('Mean Lexicon Mentions per 100 Tokens', fontsize=12, fontweight='bold')
plt.title('Bug / Money / Gameplay Language Across Negative Reviews', fontsize=14, fontweight='bold')
plt.xticks(rotation=15)
plt.legend(title='Category', fontsize=9)

```

```
plt.grid(axis='y', alpha=0.3, linestyle='--')
plt.tight_layout()
plt.savefig('plots/lexicon_subgroups.png', dpi=300, bbox_inches='tight')
plt.show()
```



In [26]: # identify which topics are most different between recommended and not recommended

```
topic_cols = [c for c in df.columns if c.startswith('topic_')]

topic_means = df.groupby('recommended')[topic_cols].mean().T
topic_means.columns = ['mean_not_recommended', 'mean_recommended']
topic_means['diff_rec_minus_not'] = topic_means['mean_recommended'] - topic_means['mean_not_recommended']
topic_means_sorted = topic_means.reindex(topic_means['diff_rec_minus_not'].sort_values(ascending=False).head(10))
display(topic_means_sorted.head(10))

top_topics_for_model = list(topic_means_sorted.head(5).index)

plt.style.use('seaborn-v0_8-whitegrid')

plt.rcParams['mathtext.default'] = 'regular'

topic_viz = topic_means_sorted.head(10).copy()
topic_viz = topic_viz.sort_values('diff_rec_minus_not', ascending=False)
```

```

fig, ax = plt.subplots(figsize=(10, 6))

colors = ['#2E86AB' if x < 0 else '#A23B72' for x in topic_viz['diff_rec_minus_not_rec']]

bars = ax.barh(topic_viz.index, topic_viz['diff_rec_minus_not_rec'], color=colors, alpha=0.8, edgecolor='black', linewidth=0.5)

ax.axvline(x=0, color='black', linestyle='-', linewidth=0.8, zorder=10)

ax.set_xlabel('Difference in Topic Proportion\n(Recommended - Not Recommended)', fontweight='bold')
ax.set_ylabel('Topic', fontsize=12, fontweight='bold')
ax.set_title('Topic Differences Between Recommended and Not Recommended')
ax.grid(axis='x', alpha=0.3, linestyle='--')
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

legend_elements = [
    Patch(facecolor='#2E86AB', alpha=0.8, label='More common'),
    Patch(facecolor='#A23B72', alpha=0.8, label='More uncommon')
]

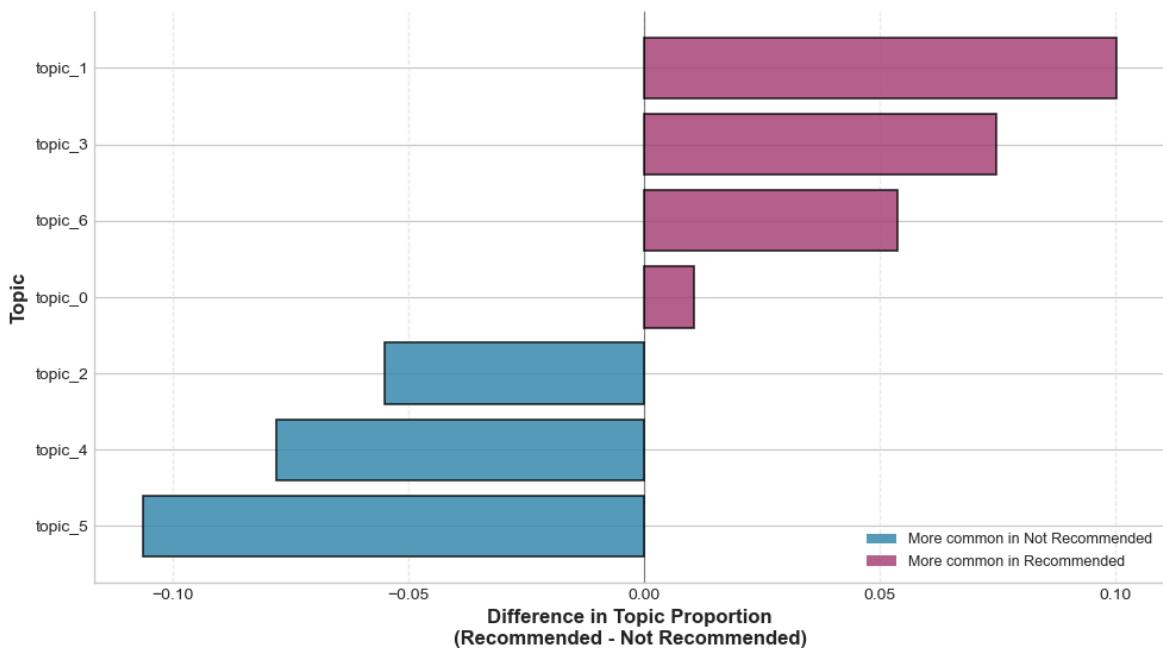
ax.legend(handles=legend_elements, loc='lower right', fontsize=10)

plt.tight_layout()
plt.savefig('plots/topic_differences.png', dpi=300, bbox_inches='tight')
plt.show()

```

	mean_not_recommended	mean_recommended	diff_rec_minus_not_rec
topic_5	0.233699	0.127372	-0.106327
topic_1	0.057900	0.158244	0.100344
topic_4	0.170016	0.091894	-0.078122
topic_3	0.179093	0.253846	0.074753
topic_2	0.127686	0.072589	-0.055107
topic_6	0.131555	0.185309	0.053754
topic_0	0.100052	0.110747	0.010695

Topic Differences Between Recommended and Not-Recommended Reviews



Step 5: Logistic Regression for Hypothesis Testing

```
In [27]: # logistic regression to find coefficients of each feature

lex_cols = ['quality_count', 'money_count', 'gameplay_count']
base_features = lex_cols + ['token_count']

topic_cols = [c for c in df.columns if c.startswith('topic_')]
selected_topics = [t for t in top_topics_for_model if t in topic_cols]

feature_cols = base_features + selected_topics
print("Features used in logistic regression:")
print(feature_cols)

model_df = df.dropna(subset=feature_cols + ['recommended'])

X = model_df[feature_cols]
y = model_df['recommended']

log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X, y)

coef_series = pd.Series(log_reg.coef_[0], index=feature_cols)
odds_ratios = np.exp(coef_series)
odds_ratios_df = pd.DataFrame({}
```

```

        'coef': coef_series,
        'odds_ratio': odds_ratios
    })

print("\nOdds ratios for lexicon predictors:")
print(odds_ratios_df.loc[lex_cols, :])

print("\nOdds ratios for selected topic predictors:")
print(odds_ratios_df.loc[selected_topics, :])

```

Features used in logistic regression:

```
['quality_count', 'money_count', 'gameplay_count', 'token_count', 'topic_5', 'topic_1', 'topic_4', 'topic_3', 'topic_2']
```

Odds ratios for lexicon predictors:

	coef	odds_ratio
quality_count	0.011841	1.011911
money_count	-0.194068	0.823602
gameplay_count	0.048951	1.050169

Odds ratios for selected topic predictors:

	coef	odds_ratio
topic_5	-1.995350	0.135966
topic_1	2.130700	8.420762
topic_4	-2.351613	0.095216
topic_3	-0.028645	0.971762
topic_2	-2.049519	0.128797

```
In [28]: or_df = odds_ratios_df.copy()
or_df = or_df.reset_index().rename(columns={'index': 'feature'})
or_df = or_df.sort_values('odds_ratio')

plt.figure(figsize=(8, 6))

plt.axvline(x=1.0, color='black', linestyle='--', linewidth=2)

plt.scatter(or_df['odds_ratio'], or_df['feature'], s=60, color='red')

plt.xscale('log')
plt.xlabel('Odds Ratio (log scale)', fontsize=12, fontweight='bold')
plt.ylabel('Feature', fontsize=12, fontweight='bold')
plt.title('Logistic Regression Odds Ratios\n(Negative reviews)')

plt.grid(axis='x', alpha=0.3, linestyle='--')
plt.tight_layout()
plt.savefig('plots/odds_ratios.png', dpi=300, bbox_inches='tight')
plt.show()
```

**Logistic Regression Odds Ratios
(Negative reviews: Recommend vs Not Recommend)**

